

AI proposes. *Rules decide*. People approve.

Word-for-word script for the talk, the live demo and the questions. Serif lines are what you say. Orange mono lines are what you do. Speaker names are a suggested split; swap them freely.

Before you go on stage

T-30 minutes, on the Guidewire VM

Start and check the services

1. Open a **cmd** console in `C:\ProvenPath` and run `vm\start-all.cmd` (it leaves anything already running alone).
2. Health: `curl http://localhost:8080/api/v1/health` shows `"llmMode":"live"` and `"packageBuilder":"provenpath.pcexport.PackageBuilder"`.
3. The **ProvenPath PC agent** console shows `ProvenPath PC agent 'vm-pcagent' started`.
4. PolicyCenter answers at `http://localhost:8180/pc`; log in as `su` / `gw` in a second tab.

Browser tabs, in this order

1. The slide deck (Present mode).
2. Mission Control landing page: `http://localhost:3000` (scrolled to the top).
3. PolicyCenter, logged in, on the Desktop.

Rehearse these three once

- **Does Gemini block?** Its instructions mention the 50% cap, so it may produce a compliant product first time. If it does, use the tamper test for the live block (script below covers both).
- **PolicyCenter refusing ₹60L on screen** has not been seen by a human yet. Try it once; if it doesn't show, skip that beat.
- **A full deploy** takes about 4 minutes of PolicyCenter restart. Know what you will show while it runs.

The demo prompt (copy it now)

Cyber insurance for a Chennai healthtech startup handling patient records: ₹1 crore aggregate limit, ₹60 lakh ransomware and extortion cover, business interruption cover kicking in after 6 hours of downtime, data breach response and privacy liability, regulatory fines where insurable by law, annual turnover ₹80 crore, minimum premium ₹40,000

1· The hook

Shaurya slide 1 · 45 s

Stand still. Don't click. Let the cover sit for a beat.

“Quick question. How long does it take an insurer to launch a new insurance product?”

Pause two seconds.

“Weeks. Coverages, limits, deductibles, exclusions, and then someone has to configure all of it inside Guidewire PolicyCenter.”

“Now ask ChatGPT or Gemini to do it. You get a full product in about ten seconds. It looks perfect. And somewhere inside it, there's a number the regulator never allowed.”

“We're ProvenPath. We let AI build insurance products at machine speed, and we make it impossible for a wrong one to reach PolicyCenter. AI proposes. Rules decide. People approve.”

“I'm Shaurya, with Chinmay, Dhriti and Vaishnavi.”

2· The problem

Shaurya slides 2-3 · 1 min 15 s

Slide 2.

“Configuring a product by hand takes around three weeks; that's the estimate in our product brief. A language model does it in seconds. Fast is great. But insurance is regulated, and a model is fluent, confident, and sometimes wrong.”

Slide 3.

“Here are three mistakes we actually caught. The model offered forty lakh of ransomware cover on a fifty lakh policy, when the cap is half the aggregate. It dropped the war exclusion that IRDAI lists as a standard exclusion. And it quoted a CERT-In direction that says incidents shall be reported within six hours, and changed shall to may.”

“One word. That's the difference between a legal duty and a suggestion. In a regulated product, sometimes wrong is too often.”

3 · Solution and architecture

Shaurya slides 4-6 · 2 min

Slide 4.

“So we built a pre-commit gate. It sits on the only road into PolicyCenter.”

“Step one, the AI proposes: Gemini turns a plain-English request into PolicyCenter coverages, and every clause has to cite the regulation it relies on. Step two, rules decide: twenty-three deterministic rules check it. One failure blocks the whole thing, and the named failure goes back to the AI to fix. Step three, a named compliance reviewer signs it. Only then does it ship.”

“If it's blocked or not approved, not a single file reaches PolicyCenter. The screen literally says: files written, zero.”

Slide 5, the architecture.

“Four parts, and trust only flows one way. Gemini only proposes; it never decides anything. Our backend, written in Gosu, Guidewire's own language, runs the gate, the review and the signing. An agent sitting right next to PolicyCenter pulls the approved package, checks the signature itself, installs it, and PolicyCenter confirms it through its own API.”

“And this is the real Guidewire PolicyCenter 10, on the VM Guidewire gave us. No mock.”

Slide 6, the pipeline. Keep it short.

“Every step is an event: proposed, verified, blocked, repaired, approved, deployed. They go into an append-only log the database won't let anyone edit, and stream live to the screen. The UI never calculates anything; it just shows what the system did.”

4 · Rules and regulations

Dhriti slides 7-8 · 1 min 45 s

Slide 7.

“The gate has six layers and twenty-three rules. Types and schema first. Then ranges: the aggregate, the extortion sublimit, deductibles, the business-interruption waiting period, and turnover. Then consistency, then mandatory covers and exclusions, then source integrity, and finally grounding: every number the AI writes in its summary must match a number we verified.”

“It's all-or-nothing, and it's deterministic. Same proposal, same rules, same verdict, same hash. Every time.”

Slide 8.

“And the rules are real. They cite IRDAI's guidance document on cyber insurance, the CERT-In directions on six-hour incident reporting, the MSME notification from March 2025 that sets the turnover limit at five hundred crore, and the DPDP Act.”

“We downloaded the official documents and store the text exactly. When the AI cites a regulation, the gate checks it byte for byte.”

“One honest detail: no Indian regulator sets a number like extortion must be under fifty percent. Insurers set those in their own product filings. So those rules cite our underwriting guideline, and the screen labels it as ours, not IRDAI's.”

“My favourite detail: the DPDP Act's breach duties only come into force around May 2027. So if the AI cites it for a policy starting now, the gate refuses it. Not in force yet.”

5 • Trust and tamper-proofing

Chinmay slide 9 • 1 min

Slide 9.

“How do we know what was verified is exactly what ends up in PolicyCenter? Every step is sealed. The proposal is hashed, the rules are hashed, the verdict is hashed. A passed run gets a signed gate token. The package gets a signature over every file and every cap.”

“And the agent next to PolicyCenter trusts nobody. Change one word of the CERT-In text: blocked. Raise a cap in the manifest: the signature fails. Swap a file: hash mismatch. Try to smuggle code in: path not allowed. Every one of those is an automated test.”

“Enough slides. Let's build one live.”

Landing page (30 s)

Switch to the Mission Control tab at `localhost:3000`. Move the mouse slowly over the hero; the proposals stream into the gate.

“This is ProvenPath. The animation is the idea: proposals flowing into a gate. Most pass. Some get stopped.”

Point at the small status line under the buttons.

“That line is live: the gate backend is online, this is the fingerprint of today's rule set, and the planner is Gemini.”

Scroll once through “How it works”, then back up. Don't read it out.

“The page explains it in four steps. Let's just launch it.”

Launch (15 s)

Click **Launch Mission Control**. The circle opens and three checks tick green.

“It really checks the backend, loads the twenty-three rules and finds the compliance reviewers before it opens.”

The request (1 min)

Paste the prompt. Say it while it's on screen, then press **Enter**.

“A Chennai healthtech startup, handling patient records. One crore of cover. Sixty lakh of ransomware cover. Business interruption after six hours. Let's see what the AI builds.”

Point at the stage rail at the top: Propose, Verify, Review, Deploy. Then at the graph as it lights up.

“Gemini is calling tools, proposing the product, adding coverages. Now the gate: six columns, one per layer, twenty-three rules lighting up as they run.”

If it blocks (expected)

The red **Blocked** stamp appears. Point to the right panel: Proposed vs Allowed.

“Blocked. Sixty lakh proposed, fifty lakh allowed: the extortion cover can't be more than half the aggregate. And six hours of waiting period, when the minimum is eight. Look here: files written to PolicyCenter, zero.”

“Now watch: nobody touches anything. The gate sends those two named failures back to Gemini, it repairs exactly those two, and the next iteration runs the whole gate again.”

Wait for the green **Passed** stamp.

“Twenty-three out of twenty-three. Now it needs a human.”

If it passes first time

The green **Passed** stamp appears straight away.

“It passed first time: Gemini already respected the caps. Good AI. But what if it had paraphrased a regulation? Let's be the bad AI.”

Open the ... menu, click **Tamper test**. The right panel shows **Blocked**, CYB-SRC-002.

“We changed one word of the real CERT-In direction, shall to may, and sent it to the same gate. Blocked at the source layer. Nothing stored, nothing deployed.”

Click the **×** on the tamper result to return to the sign-off.

Provenance (30 s)

In the Compliance review panel, click **Cyber Extortion Coverage**. The sheet slides in.

“This is what the reviewer sees. The clause, every rule that checked it, and the regulation it cites: IRDAI's own words, marked as matching the source byte for byte, with a link to the official document.”

Press **Esc**.

Approve and deploy (30 s)

Click **Approve for PolicyCenter**, then **Continue to deployment**, then **Deploy to PolicyCenter**.

“Approved by a named reviewer, logged against the verdict hash. Now deploy.”

“The backend builds a signed package. The agent next to PolicyCenter pulls it, checks the signature and every file hash itself, backs up the current files, writes the product, and restarts PolicyCenter.”

While PolicyCenter restarts (~4 min)

The timer runs under “Restarting”. Use the time; don't stand silent.

“A PolicyCenter restart takes about four minutes for a change like this; we measured it on the VM. While it restarts, two more things.”

Open ... → **Re-verify this run**. The toast shows “Same verdict, same hash”.

“Re-verify runs the whole gate again on the stored proposal. Same verdict, same hash. That's what deterministic means, and it's what an auditor needs.”

Click **Metrics** in the top bar.

“We also test the gate itself: forty labelled proposals. Zero out of twenty non-compliant ones got through, and zero out of twenty good ones were wrongly blocked.”

Close Metrics (Esc). Switch to slide 11 if there's still time; come back when the panel says Ready.

Live in PolicyCenter (30 s)

The headline turns green: **Live in PolicyCenter**. Point at the last line of the panel.

“Confirmed. That's PolicyCenter's own API telling us the product, the coverages and every capped term are there. And the caps we verified are now PolicyCenter's own limits.”

Chinmay, only if rehearsed: switch to the PolicyCenter tab, open the SMCyber submission's coverages, type 6000000 in Cyber Extortion Limit.

“So even a human underwriter who types sixty lakh into PolicyCenter tomorrow gets refused. By PolicyCenter itself.”

7 · Features, results, honesty

Vaishnavi slides 12-15 · 1 min 30 s

Slide 12, features. Don't read every card.

“Everything you saw lives in Mission Control: the live rule graph, the repair loop, provenance, sign-off, signed deploys, the tamper test and re-verify. There's also an MCP endpoint, so any AI agent can use this gate, not just ours.”

Slide 13.

“In numbers: zero false passes, zero false blocks on our labelled set, sixty-two automated tests, twenty-three rules, ten official texts quoted exactly.”

Slide 14.

“And what's not done yet, openly: the numeric caps are our own guideline, the reviewer has no login yet, and PolicyCenter doesn't have rupees as a currency, so rupee amounts sit in its dollar fields. The screen never shows a number the system didn't produce.”

Slide 15 is optional; skip it if you're behind time.

8 · The future

Dhriti slides 16-17 · 45 s

Slide 16.

“Version one is cyber. But the engine doesn't care what the product is: rules are data. Version two brings every line, each backed by real law: public liability under the Public Liability Insurance Act, fire for small businesses on IRDAI's Bharat Sookshma and Laghu products, employees' compensation, motor third party, surety bonds. Each one maps onto a line PolicyCenter already has.”

Slide 17, the team. Each person says their one line, five seconds each.

9 • Close

Shaurya slide 18 • 20 s

“AI will write insurance products. That's already happening. The question is whether you can prove every one of them is right before it goes live.”

“With ProvenPath, you can. AI proposes, rules decide, people approve. Thank you. Happy to take questions.”

Are the regulations real?

Yes. IRDAI's cyber guidance, the CERT-In directions, the MSME notification and the DPDP Act are quoted exactly from the official documents, and the gate checks citations byte for byte. The numeric caps are our own underwriting guideline, and we label them that way.

What if Gemini is wrong in a way your rules don't cover?

Then the clause has no rule that matches it, and the gate marks it NEEDS_REVIEW, which blocks. Unknown means no, not yes. And a human still signs every product.

What if Gemini is down?

The planner falls back to recorded proposals and says so in the trace, and the gate still runs live. The gate never depends on the AI.

Why Gosu?

It's Guidewire's own language, so the people who configure PolicyCenter can read and extend the gate. The backend, the gate, the package builder and the agent are all Gosu.

Why not just use PolicyCenter's own validation?

We do use it: the verified caps become PolicyCenter's own term limits. But PolicyCenter can't check that an AI's citation matches a regulation, or that a human signed the product. That's the gate's job, before anything is written.

How is this different from a guardrail library?

Guardrails filter model output. We sit on the write path to the core system, and every decision is deterministic, hashed, signed and replayable for an auditor.

Does it scale to other products?

Yes. Rules, sources and PolicyCenter templates are data per product; slide 16 is the plan, starting with public liability.

How long did it take?

Two days, four people, one Guidewire VM. GitHub was our only link to the VM because the college network blocks tunnels.

If something breaks

stay calm, say it plainly, switch

WHAT HAPPENS	WHAT TO SAY	WHAT TO DO
Trace shows fallback (Gemini failed)	“Gemini didn't answer, so the planner fell back to a recorded proposal. The gate is still live.”	Carry on; the gate, review and deploy are all real.
Backend offline on launch	“The backend is down, so let me show you a real recorded run.”	Click Watch the recorded run in the launch screen.
Deploy stuck at Queued	“The agent isn't picking it up; let me show the last deploy.”	Check the PC agent console. Otherwise stay on slide 11.
Restart runs long	“A first install compiles PolicyCenter for twenty minutes; ours is a values change.”	Go on to the results slides; come back at the end.
PolicyCenter doesn't refuse ₹60L	Say nothing about it.	Skip that beat; the ProductModelAPI confirmation already proves the install.